

```
...
```

```
[X]- find dataset (imdb sentiment analysis)
[X]- divide into train/test
[X]- create graph showing distribution of target classes
[X]- describe dataset and what the model should be able to predict
- using sklearn, try:
  [X]- naive bayes
  [X]- logistic regression
  [X]- neural networks
- write up analysis of performance of various approaches
- accuracy does not determine grade -> quality of analysis determines grade
...
```

```
import csv
import pandas as pd
# first row is review and second row is label (0 = negative, 1 = positive)
#df = pd.read_csv('movie.csv', header=0, usecols=[1,2], quoting=csv.QUOTE_NONE, enco
df = pd.read_csv('movie.csv')
print('rows and columns:', df.shape)
df.head()
```

```
↗ rows and columns: (40000, 2)
```

	text	label
0	I grew up (b. 1965) watching and loving the Th...	0
1	When I put this movie in my DVD player, and sa...	0
2	Why do people who do not know what a particula...	0
3	Even though I have great interest in Biblical ...	0
4	Im a die hard Dads Army fan and nothing will e...	1

```
# text preprocessing
import nltk
nltk.download('stopwords')

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
True
```

```
# clean data (no duplicates and NaN)
df.drop_duplicates(inplace=True)
no_of_nan_values=df.isna().sum().sum()
print(no_of_nan_values)
```

```
0
```

```
df['text_word_count']=df['text'].apply(lambda x:len(x.split()))
print(type(df['text_word_count']))
print(df['text_word_count'])
```

```
<class 'pandas.core.series.Series'>
0      151
1      326
2      184
3       69
4      178
...
39995   541
39996    50
39997   168
39998   168
39999   137
Name: text_word_count, Length: 39723, dtype: int64
```

```
# text preprocessing
from nltk.corpus import stopwords
from sklearn.feature_extraction.text import TfidfVectorizer

stopwords = set(stopwords.words('english'))
#vectorizer = TfidfVectorizer(stop_words=list(stopwords))
vectorizer_b = TfidfVectorizer(stop_words=list(stopwords), binary=True)
# .{column names}
X = df.text #features
print(type(X))
Y = df.label #targets
print(type(Y))
X.head()
```

```
<class 'pandas.core.series.Series'>
<class 'pandas.core.series.Series'>
0    I grew up (b. 1965) watching and loving the Th...
1    When I put this movie in my DVD player, and sa...
2    Why do people who do not know what a particula...
3    Even though I have great interest in Biblical ...
4    Im a die hard Dads Army fan and nothing will e...
Name: text, dtype: object
```

```
# divide into train/test
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.25, train_size=0.75, ra
X_train.shape
```

```
(29792,)
```

```
# apply tfidf vectorizer
X_train_list = X_train.values.astype('U').tolist()
X_test_list = X_test.values.astype('U').tolist()
X_train = vectorizer_b.fit_transform(X_train_list) # fit and transform the train data
X_test = vectorizer_b.transform(X_test_list) # transform only the test data


import seaborn as sb
import matplotlib.pyplot as plt

plt.figure(figsize=(10,10))
#colors=['#AB47BC','#6495ED']
colors = [sb.color_palette('pastel')[0], sb.color_palette('pastel')[4]]
plt.pie(df['label'].value_counts(),labels=['Positive','Negative'],autopct='%.1f%%',colors=col
plt.ylabel('Movie Sentiment');
```

Describe dataset and what the model should be able to predict:

This dataset contains movie reviews from IMDB and consists of two columns: text and label. The text column contains the text of the review and the label column either has a 0 for a negative label or a 1 for a positive label. This model should be able to predict the sentiment of movie reviews on IMDB.

▼ Naive Bayes

```
from sklearn.naive_bayes import MultinomialNB
#from sklearn.naive_bayes import BernoulliNB
```

```
naive_bayes = MultinomialNB()
naive_bayes.fit(X_train, Y_train)
# naive_bayes2 = BernoulliNB()
# naive_bayes2.fit(X_train, Y_train)
```

▼ MultinomialNB
MultinomialNB()

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusio
```

```
# make predictions on the test data
pred = naive_bayes.predict(X_test)
```

```
# print confusion matrix
# confusion matrix has this form
#      tp   fp
#      fn   tn
print(confusion_matrix(Y_test, pred))
```

```
[[3508  493]
 [ 625 3319]]
```

```
from sklearn.metrics import classification_report
print(classification_report(Y_test, pred))
```

	precision	recall	f1-score	support
0	0.85	0.88	0.86	4001
1	0.87	0.84	0.86	3944
accuracy			0.86	7945

macro avg	0.86	0.86	0.86	7945
weighted avg	0.86	0.86	0.86	7945

```
# print('negative size in test data:',Y_test[Y_test==0].shape[0])  
# print('test size: ', len(Y_test))  
baseline = Y_test[Y_test==0].shape[0] / Y_test.shape[0]  
print("baseline: " + str(baseline))
```

```
baseline: 0.5035871617369415
```

► Logistic Regression

[] ↳ 1 cell hidden

► Neural Network

[] ↳ 4 cells hidden

✓ 0s completed at 11:44 PM

